

Linear Regression.Project

October 26, 2016

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
%matplotlib inline
```

**** Read in the Ecommerce Customers csv file as a DataFrame called customers.***** Avg. Session Length: Average session of in-store style advice sessions. * Time on App: Average time spent on App in minutes * Time on Website: Average time spent on Website in minutes * Length of Membership: How many years the customer has been a member.

```
In [2]: customers = pd.read_csv("Ecommerce Customers")
```

```
In [3]: customers.head()
```

```
Out[3]:
```

	Email	Address	Avatar
0	mstephenson@fernandez.com	835 Frank Tunnel\nWrightmouth, MI 82180-9605	Violet
1	hduke@hotmail.com	4547 Archer Common\nDiazchester, CA 06566-8576	DarkGreen
2	pallen@yahoo.com	24645 Valerie Unions Suite 582\nCobbborough, D...	Bisque
3	riverarebecca@gmail.com	1414 David Throughway\nPort Jason, OH 22070-1220	SaddleBrown
4	mstephens@davidson-herman.com	14023 Rodriguez Passage\nPort Jacobville, PR 3...	MediumAquaMarine

	Avg. Session Length	Time on App	Time on Website	Length of Membership
0	34.497268	12.655651	39.577668	4.082621
1	31.926272	11.109461	37.268959	2.664034
2	33.000915	11.330278	37.110597	4.104543
3	34.305557	13.717514	36.721283	3.120179
4	33.330673	12.795189	37.536653	4.446308

	Yearly Amount Spent
0	587.951054
1	392.204933
2	487.547505
3	581.852344
4	599.406092

```
In [4]: customers.describe()
```

```
Out[4]:
```

	Avg. Session Length	Time on App	Time on Website
count	500.000000	500.000000	500.000000
mean	33.053194	12.052488	37.060445
std	0.992563	0.994216	1.010489
min	29.532429	8.508152	33.913847
25%	32.341822	11.388153	36.349257
50%	33.082008	11.983231	37.069367
75%	33.711985	12.753850	37.716432
max	36.139662	15.126994	40.005182

	Length of Membership	Yearly Amount Spent
count	500.000000	500.000000
mean	3.533462	499.314038
std	0.999278	79.314782
min	0.269901	256.670582
25%	2.930450	445.038277
50%	3.533975	498.887875
75%	4.126502	549.313828
max	6.922689	765.518462

```
In [279]: customers.info()
```


```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 500 entries, 0 to 499
Data columns (total 8 columns):
Email                500 non-null object
Address              500 non-null object
Avatar               500 non-null object
Avg. Session Length  500 non-null float64
Time on App          500 non-null float64
Time on Website      500 non-null float64
Length of Membership  500 non-null float64
Yearly Amount Spent  500 non-null float64
dtypes: float64(5), object(3)
memory usage: 31.3+ KB
```

0.1 Exploratory Data Analysis

```
In [5]: sns.set_palette("GnBu_d")
        sns.set_style('whitegrid')
```


```
In [281]: sns.jointplot(x='Time on Website',y='Yearly Amount Spent',data=customers)
```

```
Out[281]: <seaborn.axisgrid.JointGrid at 0x120bfcc88>
```




```
In [6]: sns.jointplot(x='Time on App',y='Yearly Amount Spent',data=customers)
```

```
Out[6]: <seaborn.axisgrid.JointGrid at 0x1f4f95f198>
```



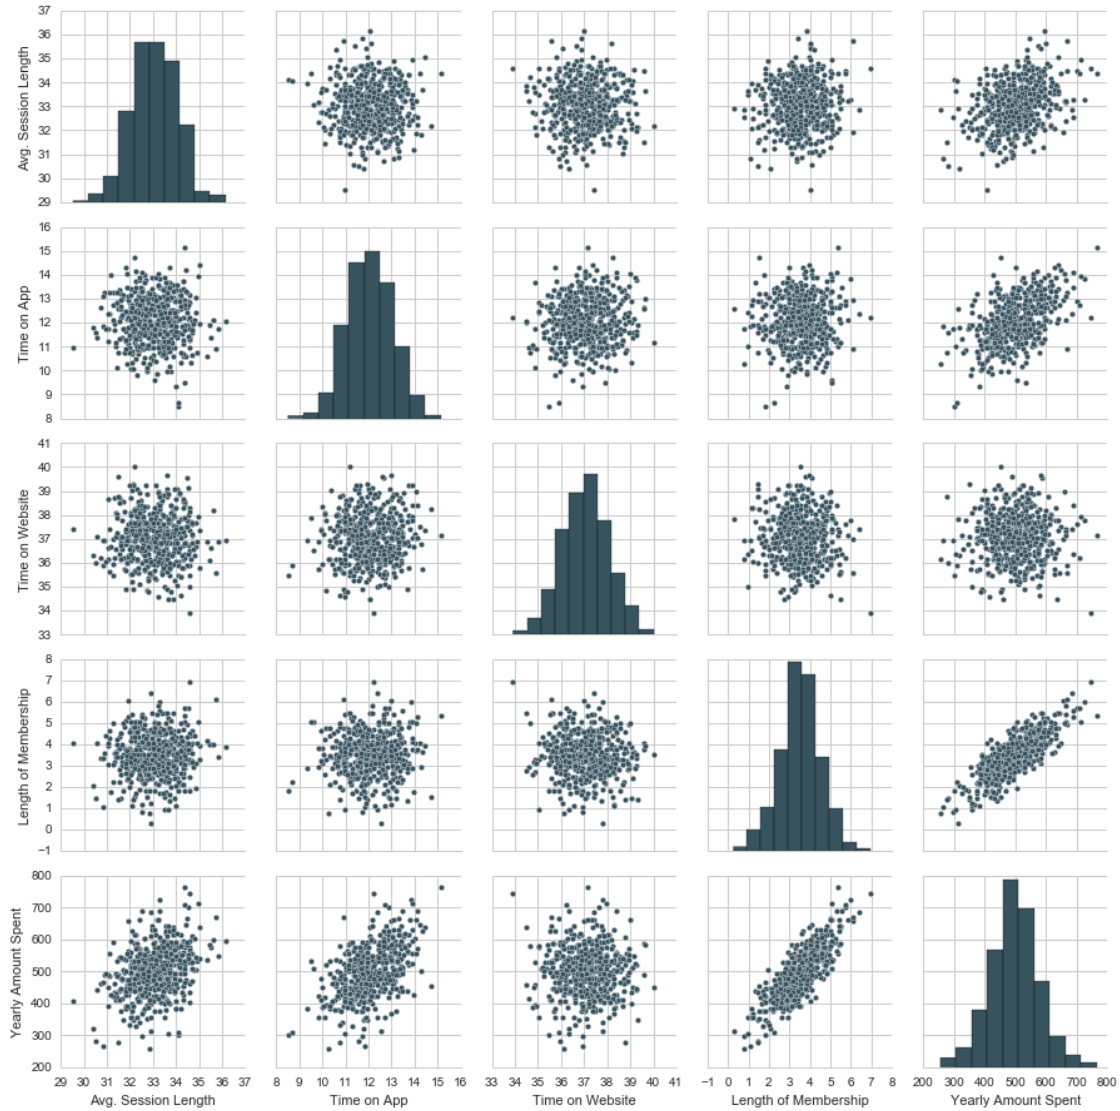
```
In [7]: sns.jointplot(x='Time on App',y='Length of Membership',kind='hex',data=customers)
```

```
Out[7]: <seaborn.axisgrid.JointGrid at 0x1f50373c88>
```




```
In [284]: sns.pairplot(customers)
```

```
Out[284]: <seaborn.axisgrid.PairGrid at 0x132fb3da0>
```



```
In [8]: sns.lmplot(x='Length of Membership',y='Yearly Amount Spent',data=customers)
```

```
Out[8]: <seaborn.axisgrid.FacetGrid at 0x1f5052aac8>
```



0.2 Training and Testing Data

```
In [9]: y = customers['Yearly Amount Spent']
```

```
In [10]: X = customers[['Avg. Session Length', 'Time on App', 'Time on Website', 'Length of Membership']]
```

**** Use cross_validation.train_test_split from sklearn to split the data into training and testing sets. Set test_size=0.3 and random_state=101****

```
In [11]: from sklearn.cross_validation import train_test_split
```

```
In [12]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=101)
```

0.3 Training the Model

```
In [13]: from sklearn.linear_model import LinearRegression
```

```
In [14]: lm = LinearRegression()
```

**** Train/fit lm on the training data.****

```
In [15]: lm.fit(X_train,y_train)
```

```
Out[15]: LinearRegression(copy_X=True, fit_intercept=True, n_jobs=1, normalize=False)
```

Print out the coefficients of the model

```
In [16]: print('Coefficients: \n', lm.coef_)
```

```
Coefficients:  
[ 25.98154972  38.59015875   0.19040528  61.27909654]
```


0.4 Predicting Test Data

```
In [17]: predictions = lm.predict( X_test)
```

0.5 Scatter plot

```
In [18]: plt.scatter(y_test,predictions)  
plt.xlabel('Y Test')  
plt.ylabel('Predicted Y')
```

```
Out[18]: <matplotlib.text.Text at 0x1f515fd0b8>
```



0.6 Evaluating the Model

```
In [19]: from sklearn import metrics  
  
print('MAE:', metrics.mean_absolute_error(y_test, predictions))  
print('MSE:', metrics.mean_squared_error(y_test, predictions))  
print('RMSE:', np.sqrt(metrics.mean_squared_error(y_test, predictions)))
```



```
MAE: 7.22814865343
MSE: 79.813051651
RMSE: 8.93381506698
```

```
In [22]: metrics.explained_variance_score(y_test,predictions)
```

```
Out[22]: 0.98907712318896057
```


```
In [23]: lm.score(X,y)
```

```
Out[23]: 0.98427271423360208
```

0.7 Residuals

Plot a histogram of the residuals and make sure it looks normally distributed. Use either seaborn distplot, or just plt.hist().

```
In [317]: sns.distplot((y_test-predictions),bins=50);
```



0.8 Conclusion

**** Recreate the dataframe below. ****

```
In [298]: coefficients = pd.DataFrame(lm.coef_,X.columns)
          coefficients.columns = ['Coefficient']
          coefficients
```

```
Out[298]:
```

	Coefficient
Avg. Session Length	25.981550
Time on App	38.590159
Time on Website	0.190405
Length of Membership	61.279097

```
In [24]: lm.score(X,y)
```

```
Out[24]: 0.98427271423360208
```

Interpreting the coefficients:

- Holding all other features fixed, a 1 unit increase in **Avg. Session Length** is associated with an **increase of 25.98 total dollars spent**.
- Holding all other features fixed, a 1 unit increase in **Time on App** is associated with an **increase of 38.59 total dollars spent**.
- Holding all other features fixed, a 1 unit increase in **Time on Website** is associated with an **increase of 0.19 total dollars spent**.
- Holding all other features fixed, a 1 unit increase in **Length of Membership** is associated with an **increase of 61.27 total dollars spent**.