

Support Vector Machines Project - Iris Data set

October 27, 2016

1 Support Vector Machines Project - Iris Data Set

```
In [17]: # The Iris Setosa
         from IPython.display import Image
         url = 'http://upload.wikimedia.org/wikipedia/commons/5/56/Kosaciec_szczecinkowaty_Iris_setosa.'
         Image(url,width=300, height=300)
```

Out[17]:



```
In [18]: # The Iris Versicolor
         from IPython.display import Image
         url = 'http://upload.wikimedia.org/wikipedia/commons/4/41/Iris_versicolor_3.jpg'
         Image(url,width=300, height=300)
```

Out[18]:



```
In [19]: # The Iris Virginica
         from IPython.display import Image
         url = 'http://upload.wikimedia.org/wikipedia/commons/9/9f/Iris_virginica.jpg'
         Image(url,width=300, height=300)
```

Out[19]:



The iris dataset contains measurements for 150 iris flowers from three different species.
The three classes in the Iris dataset:

```
Iris-setosa (n=50)  
Iris-versicolor (n=50)  
Iris-virginica (n=50)
```

The four features of the Iris dataset:

```
sepal length in cm  
sepal width in cm  
petal length in cm  
petal width in cm
```

1.1 Get the data

```
In [24]: import seaborn as sns  
         iris = sns.load_dataset('iris')
```

1.2 Exploratory Data Analysis

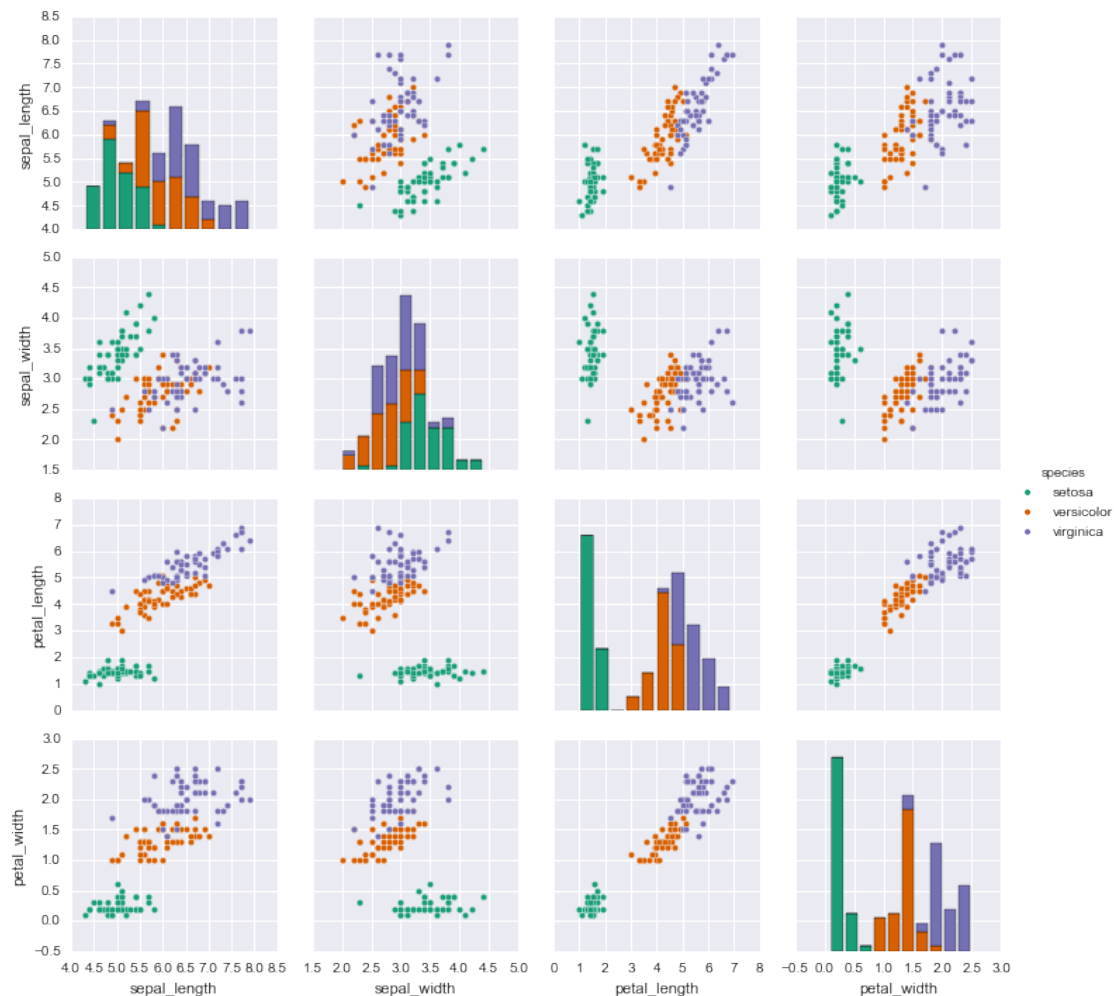
Import some libraries you think you'll need.

```
In [25]: import pandas as pd
import matplotlib.pyplot as plt
%matplotlib inline
```

**** Create a pairplot of the data set. Which flower species seems to be the most separable? ****

```
In [37]: sns.pairplot(iris,hue='species',palette='Dark2')
```


```
Out[37]: <seaborn.axisgrid.PairGrid at 0x12afb9cc0>
```



Create a kde plot of sepal_length versus sepal width for setosa species of flower.

```
In [44]: setosa = iris[iris['species']=='setosa']
sns.kdeplot( setosa['sepal_width'], setosa['sepal_length'],
            cmap="plasma", shade=True, shade_lowest=False)
```

```
Out[44]: <matplotlib.axes._subplots.AxesSubplot at 0x12f102080>
```



2 Train Test Split

**** Split your data into a training set and a testing set.****

```
In [45]: from sklearn.cross_validation import train_test_split
```

```
In [47]: X = iris.drop('species',axis=1)
         y = iris['species']
         X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.30)
```

3 Train a Model

Call the SVC() model from sklearn and fit the model to the training data.

```
In [48]: from sklearn.svm import SVC
```

```
In [49]: svc_model = SVC()
```

```
In [50]: svc_model.fit(X_train,y_train)
```

```
Out[50]: SVC(C=1.0, cache_size=200, class_weight=None, coef0=0.0,
             decision_function_shape=None, degree=3, gamma='auto', kernel='rbf',
             max_iter=-1, probability=False, random_state=None, shrinking=True,
             tol=0.001, verbose=False)
```

3.1 Model Evaluation

Now get predictions from the model and create a confusion matrix and a classification report.

```
In [51]: predictions = svc_model.predict(X_test)
```

```
In [52]: from sklearn.metrics import classification_report, confusion_matrix
```

```
In [53]: print(confusion_matrix(y_test, predictions))
```

```
[[15  0  0]
 [ 0 13  1]
 [ 0  0 16]]
```

```
In [54]: print(classification_report(y_test, predictions))
```

precision	recall	f1-score	support		
	setosa	1.00	1.00	1.00	15
	versicolor	1.00	0.93	0.96	14
	virginica	0.94	1.00	0.97	16
avg / total		0.98	0.98	0.98	45

3.2 Gridsearch Practice

```
** Import GridsearchCV from SciKit Learn.**
```

```
In [55]: from sklearn.grid_search import GridSearchCV
```

```
In [57]: param_grid = {'C': [0.1, 1, 10, 100], 'gamma': [1, 0.1, 0.01, 0.001]}
```

```
** Create a GridSearchCV object and fit it to the training data.**
```

```
In [58]: grid = GridSearchCV(SVC(), param_grid, refit=True, verbose=2)
        grid.fit(X_train, y_train)
```

Fitting 3 folds for each of 16 candidates, totalling 48 fits

```
[CV] gamma=1, C=0.1 ...
[CV] ... gamma=1, C=0.1 - 0.0s
[CV] gamma=1, C=0.1 ...
[CV] ... gamma=1, C=0.1 - 0.0s
[CV] gamma=1, C=0.1 ...
[CV] ... gamma=1, C=0.1 - 0.0s
[CV] gamma=0.1, C=0.1 ...
[CV] ... gamma=0.1, C=0.1 - 0.0s
[CV] gamma=0.1, C=0.1 ...
[CV] ... gamma=0.1, C=0.1 - 0.0s
[CV] gamma=0.1, C=0.1 ...
[CV] ... gamma=0.1, C=0.1 - 0.0s
[CV] gamma=0.01, C=0.1 ...
[CV] ... gamma=0.01, C=0.1 - 0.0s
[CV] gamma=0.01, C=0.1 ...
[CV] ... gamma=0.01, C=0.1 - 0.0s
[CV] gamma=0.01, C=0.1 ...
[CV] ... gamma=0.01, C=0.1 - 0.0s
[CV] gamma=0.001, C=0.1 ...
```

```

[CV] ... gamma=0.001, C=0.1 - 0.0s
[CV] gamma=0.001, C=0.1 ...
[CV] ... gamma=0.001, C=0.1 - 0.0s
[CV] gamma=0.001, C=0.1 ...
[CV] ... gamma=0.001, C=0.1 - 0.0s
[CV] gamma=1, C=1 ...
[CV] ... gamma=1, C=1 - 0.0s
[CV] gamma=1, C=1 ...
[CV] ... gamma=1, C=1 - 0.0s
[CV] gamma=1, C=1 ...
[CV] ... gamma=1, C=1 - 0.0s
[CV] gamma=0.1, C=1 ...
[CV] ... gamma=0.1, C=1 - 0.0s
[CV] gamma=0.1, C=1 ...
[CV] ... gamma=0.1, C=1 - 0.0s
[CV] gamma=0.1, C=1 ...
[CV] ... gamma=0.1, C=1 - 0.0s
[CV] gamma=0.01, C=1 ...
[CV] ... gamma=0.01, C=1 - 0.0s
[CV] gamma=0.01, C=1 ...
[CV] ... gamma=0.01, C=1 - 0.0s
[CV] gamma=0.01, C=1 ...
[CV] ... gamma=0.01, C=1 - 0.0s
[CV] gamma=0.001, C=1 ...
[CV] ... gamma=0.001, C=1 - 0.0s
[CV] gamma=0.001, C=1 ...
[CV] ... gamma=0.001, C=1 - 0.0s
[CV] gamma=0.001, C=1 ...
[CV] ... gamma=0.001, C=1 - 0.0s
[CV] gamma=1, C=10 ...
[CV] ... gamma=1, C=10 - 0.0s
[CV] gamma=1, C=10 ...
[CV] ... gamma=1, C=10 - 0.0s
[CV] gamma=1, C=10 ...
[CV] ... gamma=1, C=10 - 0.0s
[CV] gamma=0.1, C=10 ...
[CV] ... gamma=0.1, C=10 - 0.0s
[CV] gamma=0.1, C=10 ...
[CV] ... gamma=0.1, C=10 - 0.0s
[CV] gamma=0.1, C=10 ...
[CV] ... gamma=0.1, C=10 - 0.0s
[CV] gamma=0.01, C=10 ...
[CV] ... gamma=0.01, C=10 - 0.0s
[CV] gamma=0.01, C=10 ...
[CV] ... gamma=0.01, C=10 - 0.0s
[CV] gamma=0.01, C=10 ...
[CV] ... gamma=0.01, C=10 - 0.0s
[CV] gamma=0.001, C=10 ...
[CV] ... gamma=0.001, C=10 - 0.0s
[CV] gamma=0.001, C=10 ...
[CV] ... gamma=0.001, C=10 - 0.0s
[CV] gamma=0.001, C=10 ...
[CV] ... gamma=0.001, C=10 - 0.0s
[CV] gamma=1, C=100 ...

```



```

[CV] ... gamma=1, C=100 - 0.0s
[CV] gamma=1, C=100 ...
[CV] ... gamma=1, C=100 - 0.0s
[CV] gamma=1, C=100 ...
[CV] ... gamma=1, C=100 - 0.0s
[CV] gamma=0.1, C=100 ...
[CV] ... gamma=0.1, C=100 - 0.0s
[CV] gamma=0.1, C=100 ...
[CV] ... gamma=0.1, C=100 - 0.0s
[CV] gamma=0.1, C=100 ...
[CV] ... gamma=0.1, C=100 - 0.0s
[CV] gamma=0.01, C=100 ...
[CV] ... gamma=0.01, C=100 - 0.0s
[CV] gamma=0.01, C=100 ...
[CV] ... gamma=0.01, C=100 - 0.0s
[CV] gamma=0.01, C=100 ...
[CV] ... gamma=0.01, C=100 - 0.0s
[CV] gamma=0.001, C=100 ...
[CV] ... gamma=0.001, C=100 - 0.0s
[CV] gamma=0.001, C=100 ...
[CV] ... gamma=0.001, C=100 - 0.0s
[CV] gamma=0.001, C=100 ...
[CV] ... gamma=0.001, C=100 - 0.0s

[Parallel(n_jobs=1)]: Done 40 tasks      | elapsed: 0.2s
[Parallel(n_jobs=1)]: Done 48 out of 48 | elapsed: 0.2s finished

```

```

Out[58]: GridSearchCV(cv=None, error_score='raise',
                      estimator=SVC(C=1.0, cache_size=200, class_weight=None, coef0=0.0,
                      decision_function_shape=None, degree=3, gamma='auto', kernel='rbf',
                      max_iter=-1, probability=False, random_state=None, shrinking=True,
                      tol=0.001, verbose=False),
                      fit_params={}, iid=True, n_jobs=1,
                      param_grid={'gamma': [1, 0.1, 0.01, 0.001], 'C': [0.1, 1, 10, 100]},
                      pre_dispatch='2*n_jobs', refit=True, scoring=None, verbose=2)

```

```

In [59]: grid_predictions = grid.predict(X_test)

```

```

In [60]: print(confusion_matrix(y_test,grid_predictions))

```

```

[[15  0  0]
 [ 0 13  1]
 [ 0  0 16]]

```

```

In [61]: print(classification_report(y_test,grid_predictions))

```

precision	recall	f1-score	support		
	setosa	1.00	1.00	1.00	15
	versicolor	1.00	0.93	0.96	14
	virginica	0.94	1.00	0.97	16
avg / total		0.98	0.98	0.98	45