

Support Vector Machines project I - Cancer Data set

October 27, 2016

1 Support Vector Machines with Python

1.1 Import Libraries

```
In [51]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
%matplotlib inline
```

1.2 Get the Data

```
In [52]: from sklearn.datasets import load_breast_cancer
```

```
In [54]: cancer = load_breast_cancer()
```

The data set is presented in a dictionary form:

```
In [55]: cancer.keys()
```

```
Out[55]: dict_keys(['DESCR', 'target', 'data', 'target_names', 'feature_names'])
```

```
In [56]: cancer['feature_names']
```

```
Out[56]: array(['mean radius', 'mean texture', 'mean perimeter', 'mean area',
               'mean smoothness', 'mean compactness', 'mean concavity',
               'mean concave points', 'mean symmetry', 'mean fractal dimension',
               'radius error', 'texture error', 'perimeter error', 'area error',
               'smoothness error', 'compactness error', 'concavity error',
               'concave points error', 'symmetry error', 'fractal dimension error',
               'worst radius', 'worst texture', 'worst perimeter', 'worst area',
               'worst smoothness', 'worst compactness', 'worst concavity',
               'worst concave points', 'worst symmetry', 'worst fractal dimension'],
              dtype='<U23')
```

1.3 Set up DataFrame

```
In [12]: df_feat = pd.DataFrame(cancer['data'], columns=cancer['feature_names'])
df_feat.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 569 entries, 0 to 568
Data columns (total 30 columns):
mean radius           569 non-null float64
mean texture          569 non-null float64
```

mean perimeter	569	non-null	float64
mean area	569	non-null	float64
mean smoothness	569	non-null	float64
mean compactness	569	non-null	float64
mean concavity	569	non-null	float64
mean concave points	569	non-null	float64
mean symmetry	569	non-null	float64
mean fractal dimension	569	non-null	float64
radius error	569	non-null	float64
texture error	569	non-null	float64
perimeter error	569	non-null	float64
area error	569	non-null	float64
smoothness error	569	non-null	float64
compactness error	569	non-null	float64
concavity error	569	non-null	float64
concave points error	569	non-null	float64
symmetry error	569	non-null	float64
fractal dimension error	569	non-null	float64
worst radius	569	non-null	float64
worst texture	569	non-null	float64
worst perimeter	569	non-null	float64
worst area	569	non-null	float64
worst smoothness	569	non-null	float64
worst compactness	569	non-null	float64
worst concavity	569	non-null	float64
worst concave points	569	non-null	float64
worst symmetry	569	non-null	float64
worst fractal dimension	569	non-null	float64
dtypes: float64(30)			
memory usage: 133.4 KB			

[illegible]

```
1, 1, 1, 0, 1, 1, 0, 1, 0, 1, 0, 0, 1, 1, 1, 0, 1, 1, 1, 1, 1, 1, 1,
1, 1, 1, 1, 0, 1, 0, 0, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1,
1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 0, 0, 0, 0, 0, 0, 0, 1])
```

```
In [16]: df_target = pd.DataFrame(cancer['target'],columns=['Cancer'])
```

Now let's actually check out the dataframe!

```
In [8]: df.head()
```

```
Out[8]:
```

	mean radius	mean texture	mean perimeter	mean area	mean smoothness	\
0	17.99	10.38	122.80	1001.0	0.11840	
1	20.57	17.77	132.90	1326.0	0.08474	
2	19.69	21.25	130.00	1203.0	0.10960	
3	11.42	20.38	77.58	386.1	0.14250	
4	20.29	14.34	135.10	1297.0	0.10030	

	mean compactness	mean concavity	mean concave points	mean symmetry	\
0	0.27760	0.3001	0.14710	0.2419	
1	0.07864	0.0869	0.07017	0.1812	
2	0.15990	0.1974	0.12790	0.2069	
3	0.28390	0.2414	0.10520	0.2597	
4	0.13280	0.1980	0.10430	0.1809	

	mean fractal dimension	...	worst radius	\
0	0.07871	...	25.38	
1	0.05667	...	24.99	
2	0.05999	...	23.57	
3	0.09744	...	14.91	
4	0.05883	...	22.54	

	worst texture	worst perimeter	worst area	worst smoothness	\
0	17.33	184.60	2019.0	0.1622	
1	23.41	158.80	1956.0	0.1238	
2	25.53	152.50	1709.0	0.1444	
3	26.50	98.87	567.7	0.2098	
4	16.67	152.20	1575.0	0.1374	

	worst compactness	worst concavity	worst concave points	worst symmetry	\
0	0.6656	0.7119	0.2654	0.4601	
1	0.1866	0.2416	0.1860	0.2750	
2	0.4245	0.4504	0.2430	0.3613	
3	0.8663	0.6869	0.2575	0.6638	
4	0.2050	0.4000	0.1625	0.2364	

	worst fractal dimension
0	0.11890
1	0.08902
2	0.08758
3	0.17300
4	0.07678

[5 rows x 30 columns]

1.4 Train Test Split

```
In [57]: from sklearn.cross_validation import train_test_split
```

```
In [58]: X_train, X_test, y_train, y_test = train_test_split(df_feat, np.ravel(df_target), test_size=0.1)
```

2 Train the Support Vector Classifier

```
In [59]: from sklearn.svm import SVC
```

```
In [60]: model = SVC()
```

```
In [61]: model.fit(X_train,y_train)
```

```
Out[61]: SVC(C=1.0, cache_size=200, class_weight=None, coef0=0.0,
            decision_function_shape=None, degree=3, gamma='auto', kernel='rbf',
            max_iter=-1, probability=False, random_state=None, shrinking=True,
            tol=0.001, verbose=False)
```

2.1 Predictions and Evaluations

```
In [27]: predictions = model.predict(X_test)
```

```
In [45]: from sklearn.metrics import classification_report,confusion_matrix
```

```
In [46]: print(confusion_matrix(y_test,predictions))
```

```
[[ 0 66]
 [ 0 105]]
```

```
In [62]: print(classification_report(y_test,predictions))
```

precision	recall	f1-score	support		
	0	0.00	0.00	0.00	66
	1	0.61	1.00	0.76	105
avg / total		0.38	0.61	0.47	171

```
/Users/marci/anaconda/lib/python3.5/site-packages/sklearn/metrics/classification.py:1074: UndefinedMetricWarning: Precision is undefined because there is no true label in the data.
'precision', 'predicted', average, warn_for)
```

We can search for parameters using a GridSearch!

3 Gridsearch

```
In [63]: param_grid = {'C': [0.1,1, 10, 100, 1000], 'gamma': [1,0.1,0.01,0.001,0.0001], 'kernel': ['rbf']}
```

```
In [64]: from sklearn.grid_search import GridSearchCV
```

```
In [65]: grid = GridSearchCV(SVC(),param_grid,refit=True,verbose=3)
```

```
In [40]: grid.fit(X_train,y_train)
```

Fitting 3 folds for each of 25 candidates, totalling 75 fits

```

[CV] gamma=1, C=0.1, kernel=rbf ...
[CV] ... gamma=1, C=0.1, kernel=rbf, score=0.631579 - 0.0s
[CV] gamma=1, C=0.1, kernel=rbf ...
[CV] ... gamma=1, C=0.1, kernel=rbf, score=0.631579 - 0.0s
[CV] gamma=1, C=0.1, kernel=rbf ...
[CV] ... gamma=1, C=0.1, kernel=rbf, score=0.636364 - 0.0s
[CV] gamma=0.1, C=0.1, kernel=rbf ...
[CV] ... gamma=0.1, C=0.1, kernel=rbf, score=0.631579 - 0.0s
[CV] gamma=0.1, C=0.1, kernel=rbf ...
[CV] ... gamma=0.1, C=0.1, kernel=rbf, score=0.631579 - 0.0s
[CV] gamma=0.1, C=0.1, kernel=rbf ...
[CV] ... gamma=0.1, C=0.1, kernel=rbf, score=0.636364 - 0.0s
[CV] gamma=0.01, C=0.1, kernel=rbf ...
[CV] ... gamma=0.01, C=0.1, kernel=rbf, score=0.631579 - 0.0s
[CV] gamma=0.01, C=0.1, kernel=rbf ...
[CV] ... gamma=0.01, C=0.1, kernel=rbf, score=0.631579 - 0.0s
[CV] gamma=0.01, C=0.1, kernel=rbf ...
[CV] ... gamma=0.01, C=0.1, kernel=rbf, score=0.636364 - 0.0s
[CV] gamma=0.001, C=0.1, kernel=rbf ...
[CV] ... gamma=0.001, C=0.1, kernel=rbf, score=0.631579 - 0.0s
[CV] gamma=0.001, C=0.1, kernel=rbf ...
[CV] ... gamma=0.001, C=0.1, kernel=rbf, score=0.631579 - 0.0s
[CV] gamma=0.001, C=0.1, kernel=rbf ...
[CV] ... gamma=0.001, C=0.1, kernel=rbf, score=0.636364 - 0.0s
[CV] gamma=0.0001, C=0.1, kernel=rbf ...
[CV] ... gamma=0.0001, C=0.1, kernel=rbf, score=0.902256 - 0.0s
[CV] gamma=0.0001, C=0.1, kernel=rbf ...
[CV] ... gamma=0.0001, C=0.1, kernel=rbf, score=0.962406 - 0.0s
[CV] gamma=0.0001, C=0.1, kernel=rbf ...
[CV] ... gamma=0.0001, C=0.1, kernel=rbf, score=0.916667 - 0.0s
[CV] gamma=1, C=1, kernel=rbf ...
[CV] ... gamma=1, C=1, kernel=rbf, score=0.631579 - 0.0s
[CV] gamma=1, C=1, kernel=rbf ...
[CV] ... gamma=1, C=1, kernel=rbf, score=0.631579 - 0.0s
[CV] gamma=1, C=1, kernel=rbf ...
[CV] ... gamma=1, C=1, kernel=rbf, score=0.636364 - 0.0s
[CV] gamma=0.1, C=1, kernel=rbf ...
[CV] ... gamma=0.1, C=1, kernel=rbf, score=0.631579 - 0.0s
[CV] gamma=0.1, C=1, kernel=rbf ...
[CV] ... gamma=0.1, C=1, kernel=rbf, score=0.631579 - 0.0s
[CV] gamma=0.1, C=1, kernel=rbf ...
[CV] ... gamma=0.1, C=1, kernel=rbf, score=0.636364 - 0.0s
[CV] gamma=0.01, C=1, kernel=rbf ...
[CV] ... gamma=0.01, C=1, kernel=rbf, score=0.631579 - 0.0s
[CV] gamma=0.01, C=1, kernel=rbf ...
[CV] ... gamma=0.01, C=1, kernel=rbf, score=0.631579 - 0.0s
[CV] gamma=0.01, C=1, kernel=rbf ...
[CV] ... gamma=0.01, C=1, kernel=rbf, score=0.636364 - 0.0s
[CV] gamma=0.001, C=1, kernel=rbf ...
[CV] ... gamma=0.001, C=1, kernel=rbf, score=0.902256 - 0.0s
[CV] gamma=0.001, C=1, kernel=rbf ...
[CV] ... gamma=0.001, C=1, kernel=rbf, score=0.939850 - 0.0s
[CV] gamma=0.001, C=1, kernel=rbf ...

```

```

[CV] ... gamma=0.001, C=1, kernel=rbf, score=0.954545 - 0.0s
[CV] gamma=0.0001, C=1, kernel=rbf ...
[CV] ... gamma=0.0001, C=1, kernel=rbf, score=0.939850 - 0.0s
[CV] gamma=0.0001, C=1, kernel=rbf ...
[CV] ... gamma=0.0001, C=1, kernel=rbf, score=0.969925 - 0.0s
[CV] gamma=0.0001, C=1, kernel=rbf ...
[CV] ... gamma=0.0001, C=1, kernel=rbf, score=0.946970 - 0.0s
[CV] gamma=1, C=10, kernel=rbf ...
[CV] ... gamma=1, C=10, kernel=rbf, score=0.631579 - 0.0s
[CV] gamma=1, C=10, kernel=rbf ...
[CV] ... gamma=1, C=10, kernel=rbf, score=0.631579 - 0.0s
[CV] gamma=1, C=10, kernel=rbf ...
[CV] ... gamma=1, C=10, kernel=rbf, score=0.636364 - 0.0s
[CV] gamma=0.1, C=10, kernel=rbf ...
[CV] ... gamma=0.1, C=10, kernel=rbf, score=0.631579 - 0.0s
[CV] gamma=0.1, C=10, kernel=rbf ...
[CV] ... gamma=0.1, C=10, kernel=rbf, score=0.631579 - 0.0s
[CV] gamma=0.1, C=10, kernel=rbf ...
[CV] ... gamma=0.1, C=10, kernel=rbf, score=0.636364 - 0.0s
[CV] gamma=0.01, C=10, kernel=rbf ...
[CV] ... gamma=0.01, C=10, kernel=rbf, score=0.631579 - 0.0s
[CV] gamma=0.01, C=10, kernel=rbf ...
[CV] ... gamma=0.01, C=10, kernel=rbf, score=0.631579 - 0.0s
[CV] gamma=0.01, C=10, kernel=rbf ...
[CV] ... gamma=0.01, C=10, kernel=rbf, score=0.636364 - 0.0s
[CV] gamma=0.001, C=10, kernel=rbf ...
[CV] ... gamma=0.001, C=10, kernel=rbf, score=0.894737 - 0.0s
[CV] gamma=0.001, C=10, kernel=rbf ...
[CV] ... gamma=0.001, C=10, kernel=rbf, score=0.932331 - 0.0s
[CV] gamma=0.001, C=10, kernel=rbf ...
[CV] ... gamma=0.001, C=10, kernel=rbf, score=0.916667 - 0.0s
[CV] gamma=0.0001, C=10, kernel=rbf ...
[CV] ... gamma=0.0001, C=10, kernel=rbf, score=0.932331 - 0.0s
[CV] gamma=0.0001, C=10, kernel=rbf ...
[CV] ... gamma=0.0001, C=10, kernel=rbf, score=0.969925 - 0.0s
[CV] gamma=0.0001, C=10, kernel=rbf ...
[CV] ... gamma=0.0001, C=10, kernel=rbf, score=0.962121 - 0.0s
[CV] gamma=1, C=100, kernel=rbf ...
[CV] ... gamma=1, C=100, kernel=rbf, score=0.631579 - 0.0s
[CV] gamma=1, C=100, kernel=rbf ...
[CV] ... gamma=1, C=100, kernel=rbf, score=0.631579 - 0.0s
[CV] gamma=1, C=100, kernel=rbf ...
[CV] ... gamma=1, C=100, kernel=rbf, score=0.636364 - 0.0s
[CV] gamma=0.1, C=100, kernel=rbf ...
[CV] ... gamma=0.1, C=100, kernel=rbf, score=0.631579 - 0.0s
[CV] gamma=0.1, C=100, kernel=rbf ...
[CV] ... gamma=0.1, C=100, kernel=rbf, score=0.631579 - 0.0s
[CV] gamma=0.1, C=100, kernel=rbf ...
[CV] ... gamma=0.1, C=100, kernel=rbf, score=0.636364 - 0.0s
[CV] gamma=0.01, C=100, kernel=rbf ...
[CV] ... gamma=0.01, C=100, kernel=rbf, score=0.631579 - 0.0s
[CV] gamma=0.01, C=100, kernel=rbf ...
[CV] ... gamma=0.01, C=100, kernel=rbf, score=0.631579 - 0.0s
[CV] gamma=0.01, C=100, kernel=rbf ...

```

```

[CV] ... gamma=0.01, C=100, kernel=rbf, score=0.636364 - 0.0s
[CV] gamma=0.001, C=100, kernel=rbf ...
[CV] ... gamma=0.001, C=100, kernel=rbf, score=0.894737 - 0.0s
[CV] gamma=0.001, C=100, kernel=rbf ...
[CV] ... gamma=0.001, C=100, kernel=rbf, score=0.932331 - 0.0s
[CV] gamma=0.001, C=100, kernel=rbf ...
[CV] ... gamma=0.001, C=100, kernel=rbf, score=0.916667 - 0.0s
[CV] gamma=0.0001, C=100, kernel=rbf ...
[CV] ... gamma=0.0001, C=100, kernel=rbf, score=0.917293 - 0.0s
[CV] gamma=0.0001, C=100, kernel=rbf ...
[CV] ... gamma=0.0001, C=100, kernel=rbf, score=0.977444 - 0.0s
[CV] gamma=0.0001, C=100, kernel=rbf ...
[CV] ... gamma=0.0001, C=100, kernel=rbf, score=0.939394 - 0.0s
[CV] gamma=1, C=1000, kernel=rbf ...
[CV] ... gamma=1, C=1000, kernel=rbf, score=0.631579 - 0.0s
[CV] gamma=1, C=1000, kernel=rbf ...
[CV] ... gamma=1, C=1000, kernel=rbf, score=0.631579 - 0.0s
[CV] gamma=1, C=1000, kernel=rbf ...
[CV] ... gamma=1, C=1000, kernel=rbf, score=0.636364 - 0.0s
[CV] gamma=0.1, C=1000, kernel=rbf ...
[CV] ... gamma=0.1, C=1000, kernel=rbf, score=0.631579 - 0.0s
[CV] gamma=0.1, C=1000, kernel=rbf ...
[CV] ... gamma=0.1, C=1000, kernel=rbf, score=0.631579 - 0.0s
[CV] gamma=0.1, C=1000, kernel=rbf ...
[CV] ... gamma=0.1, C=1000, kernel=rbf, score=0.636364 - 0.0s
[CV] gamma=0.01, C=1000, kernel=rbf ...
[CV] ... gamma=0.01, C=1000, kernel=rbf, score=0.631579 - 0.0s
[CV] gamma=0.01, C=1000, kernel=rbf ...
[CV] ... gamma=0.01, C=1000, kernel=rbf, score=0.631579 - 0.0s
[CV] gamma=0.01, C=1000, kernel=rbf ...
[CV] ... gamma=0.01, C=1000, kernel=rbf, score=0.636364 - 0.0s
[CV] gamma=0.001, C=1000, kernel=rbf ...
[CV] ... gamma=0.001, C=1000, kernel=rbf, score=0.894737 - 0.0s
[CV] gamma=0.001, C=1000, kernel=rbf ...
[CV] ... gamma=0.001, C=1000, kernel=rbf, score=0.932331 - 0.0s
[CV] gamma=0.001, C=1000, kernel=rbf ...
[CV] ... gamma=0.001, C=1000, kernel=rbf, score=0.916667 - 0.0s

[Parallel(n_jobs=1)]: Done 31 tasks      | elapsed: 0.3s
[Parallel(n_jobs=1)]: Done 75 out of 75 | elapsed: 0.8s finished

```

```

[CV] gamma=0.0001, C=1000, kernel=rbf ...
[CV] ... gamma=0.0001, C=1000, kernel=rbf, score=0.909774 - 0.0s
[CV] gamma=0.0001, C=1000, kernel=rbf ...
[CV] ... gamma=0.0001, C=1000, kernel=rbf, score=0.969925 - 0.0s
[CV] gamma=0.0001, C=1000, kernel=rbf ...
[CV] ... gamma=0.0001, C=1000, kernel=rbf, score=0.931818 - 0.0s

```

```

Out[40]: GridSearchCV(cv=None, error_score='raise',
                    estimator=SVC(C=1.0, cache_size=200, class_weight=None, coef0=0.0,
                    decision_function_shape=None, degree=3, gamma='auto', kernel='rbf',
                    max_iter=-1, probability=False, random_state=None, shrinking=True,
                    tol=0.001, verbose=False),
                    fit_params={}, iid=True, n_jobs=1,
                    param_grid={'gamma': [1, 0.1, 0.01, 0.001, 0.0001], 'C': [0.1, 1, 10, 100, 1000], 'kerne
                    pre_dispatch='2*n_jobs', refit=True, scoring=None, verbose=3)

```

Inspect the best parameters found by GridSearchCV in the `best_params_` attribute, and the best estimator in the `best_estimator_` attribute:

```
In [41]: grid.best_params_
```

```
Out[41]: {'C': 10, 'gamma': 0.0001, 'kernel': 'rbf'}
```

```
In [ ]: grid.best_estimator_
```

Then you can re-run predictions on this grid object just like you would with a normal model.

```
In [48]: grid_predictions = grid.predict(X_test)
```

```
In [49]: print(confusion_matrix(y_test,grid_predictions))
```

```
[[ 60   6]
 [  3 102]]
```

```
In [50]: print(classification_report(y_test,grid_predictions))
```

precision	recall	f1-score	support		
	0	0.95	0.91	0.93	66
	1	0.94	0.97	0.96	105
avg / total		0.95	0.95	0.95	171